

Marie Delavergne

Cheops, a service mesh to geo-distribute (micro-)service applications at the Edge

Supervisor: Adrien LEBRE - Professor, IMT Atlantique

Co-supervisor: Ronan-Alexandre CHERRUEAU (2019-2020) - DGFIP

Reviewers :

Noël DE PALMA - Professor, Université Grenoble Alpes

Pierre SENS - Professor, Sorbonne Université

Examiners :

Sara BOUCHENAK - Professor, INSA Lyon

Thomas LEDOUX - Professor, IMT Atlantique



Research question

Is it possible to use applications developed for the Cloud on Edge infrastructures without changing their code?

- I Context: Cloud Computing
- II State-of-the-art
- III The approach: Cheops and scope-lang
- IV Conclusion & perspectives

Context: Cloud Computing

A simplified view of the Cloud



[Source](#)

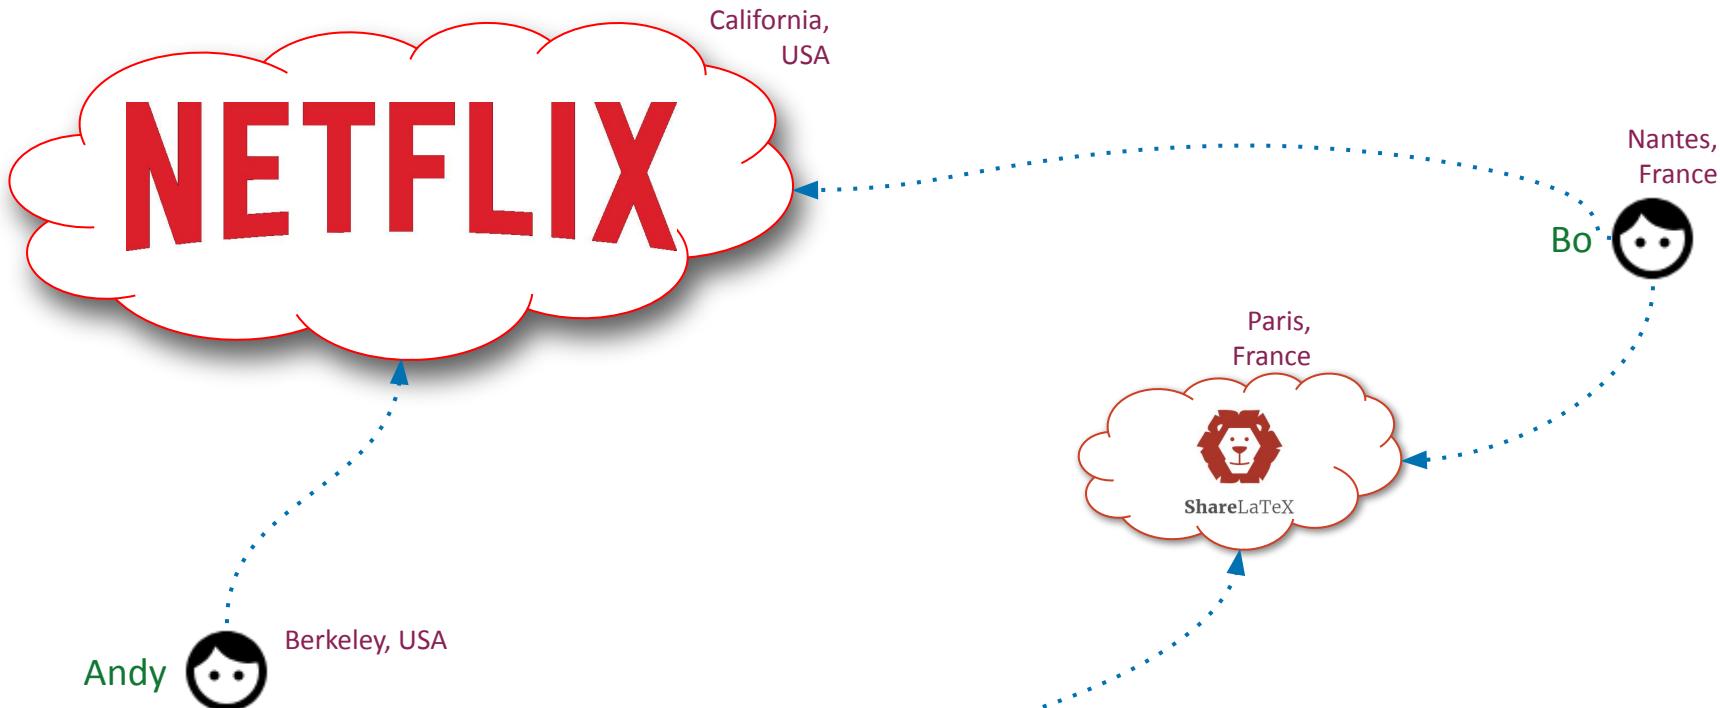


NETFLIX



ShareLaTeX

A simplified view of the Cloud

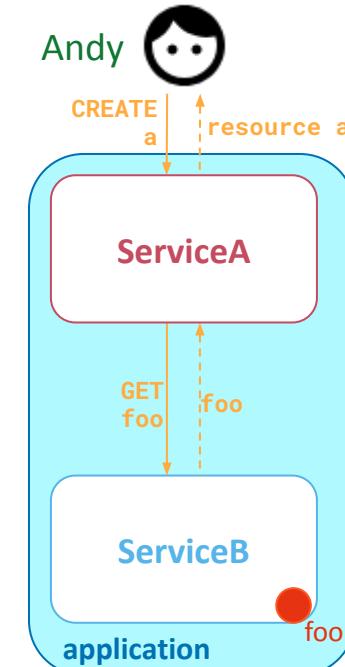


Cloud applications are (often) made of services

Services manage a subset of resources of the applications.

To create a **resourceA** on **ServiceA** using a *sub-resource* **foo** from **ServiceB**, Andy needs to give a command like this:

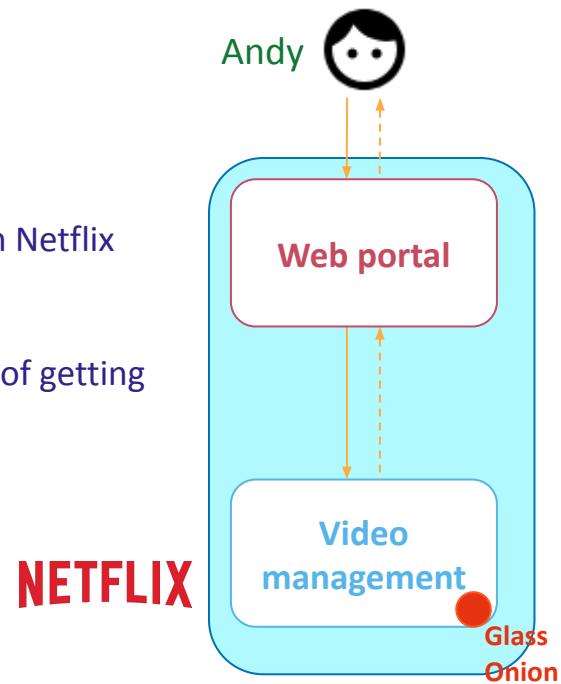
```
application create resourceA --sub-resourceB foo
```



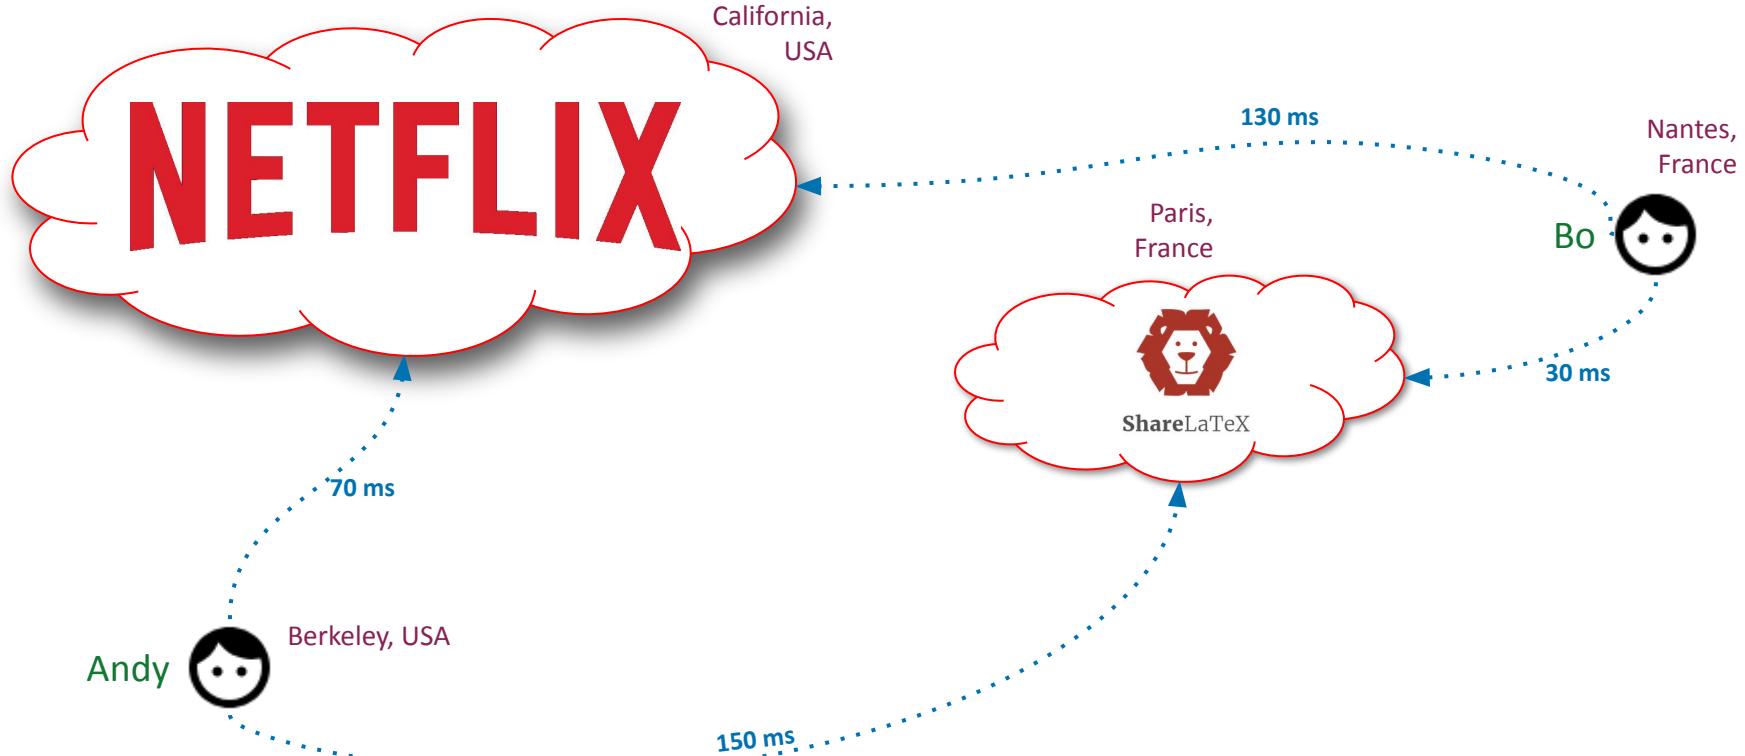
Cloud applications are (often) made of services

Andy wants to watch Glass Onion on Netflix

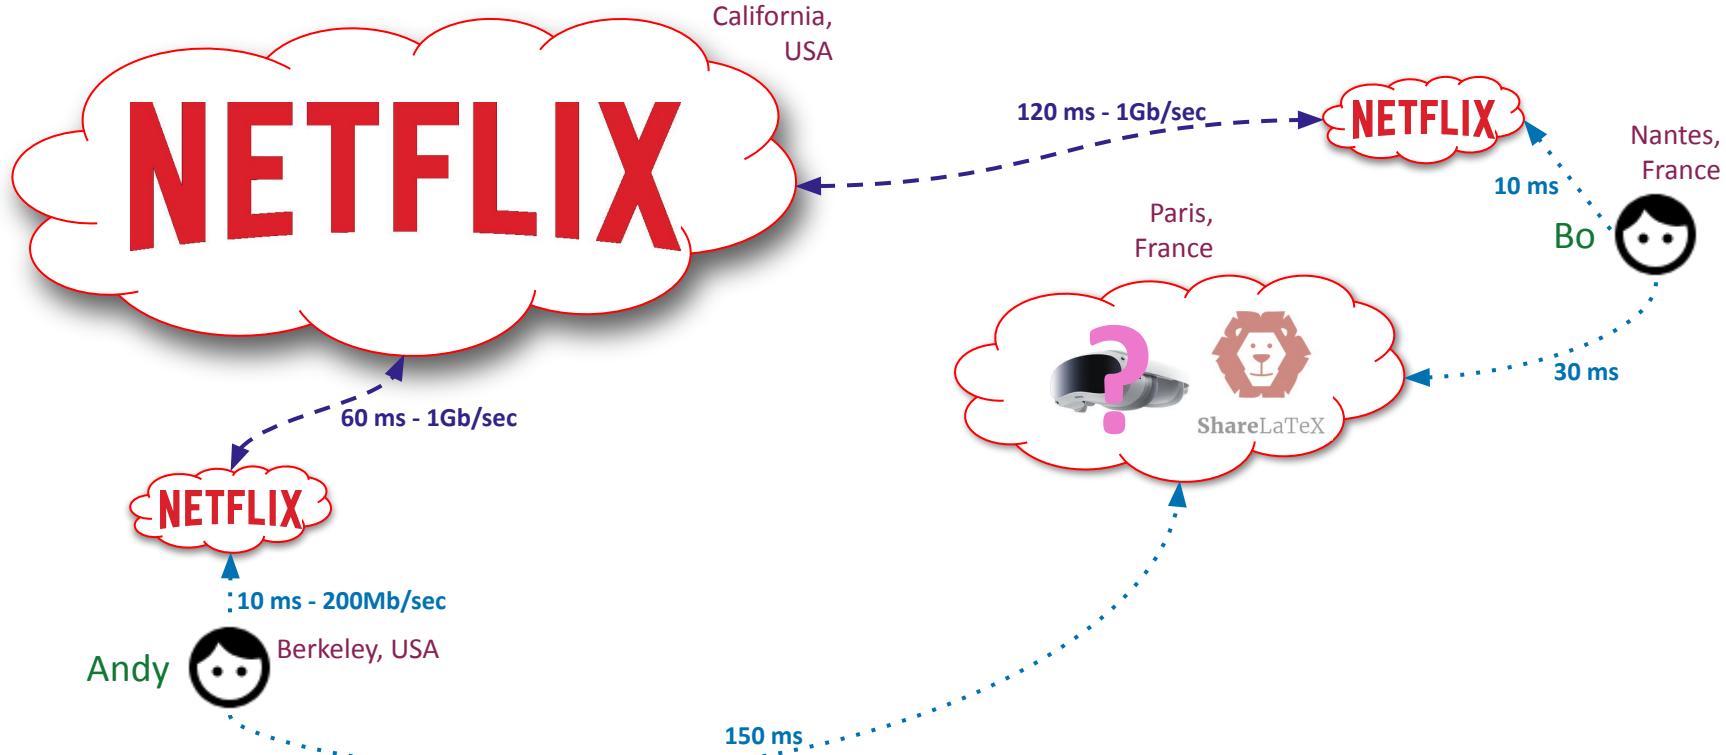
She reaches the web portal, in charge of getting the video



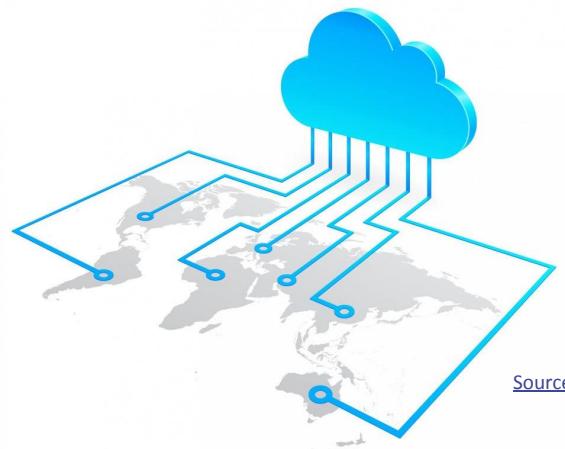
Cloud vs latency/throughput



Cloud vs latency/throughput

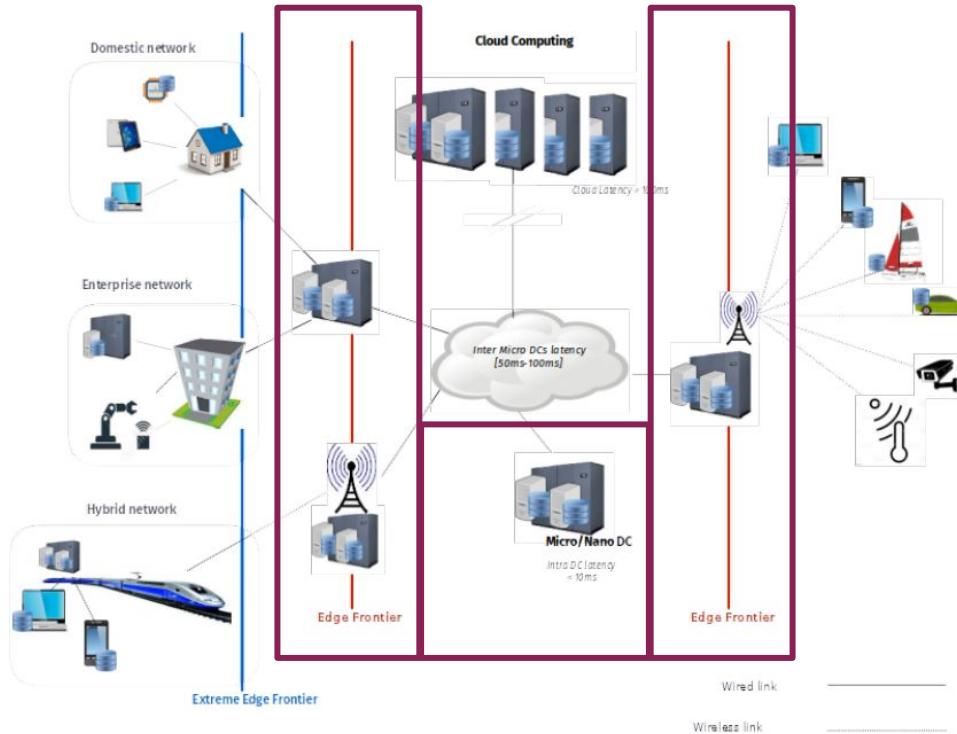


Edge Computing: What is the purpose?



Reduce the latency between the Cloud clients and the datacenters
for **latency-sensitive** and **data-intensive** applications

Many Edge definitions: our vision



A geo-distributed Cloud:
An infrastructure of micro DCs

- Cloud icon: Closer to users
- Cloud icon: Managed by one operator

Challenges:

- Cloud icon: High latency between far sites
- Cloud icon: Disconnections between sites

How can we geo-distribute an application?

NETFLIX is natively geo-distributed

Two approaches in the academic literature:

-  **ShareLatex^[1]**: efforts to geo-distribute services
-  **OpenStack^[2]**: change the DB to a geo-distributed one

We do not want to change their code,
nor put solutions to push the applications in the business code

[1] Tato et al., *ShareLatex on the Edge [...]*, ME'CC18 (ACM Middleware workshops) <https://dl.acm.org/doi/10.1145/3286685.3286687>

[2] Lebre et al., *Revising OpenStack to Operate Fog/Edge Computing infrastructures*, IC2E 2017 <https://hal.inria.fr/hal-01273427>

A geo-distributed version of ShareLaTeX?

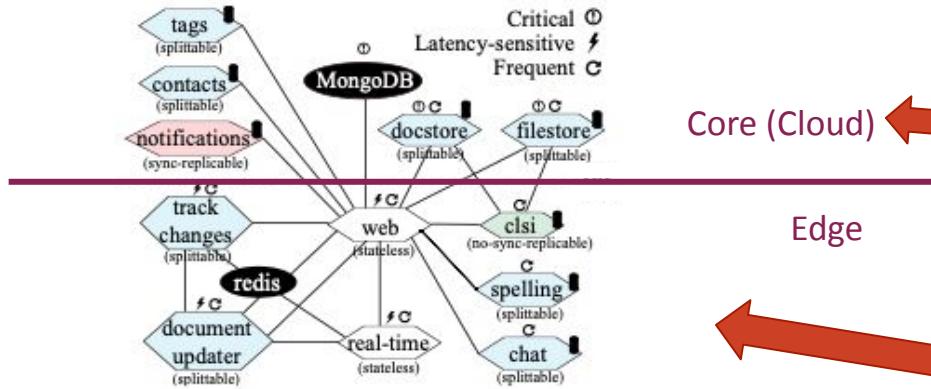


Figure 3: Our split and placement. Critical services in the core, latency-sensitive and frequent services at the edge.

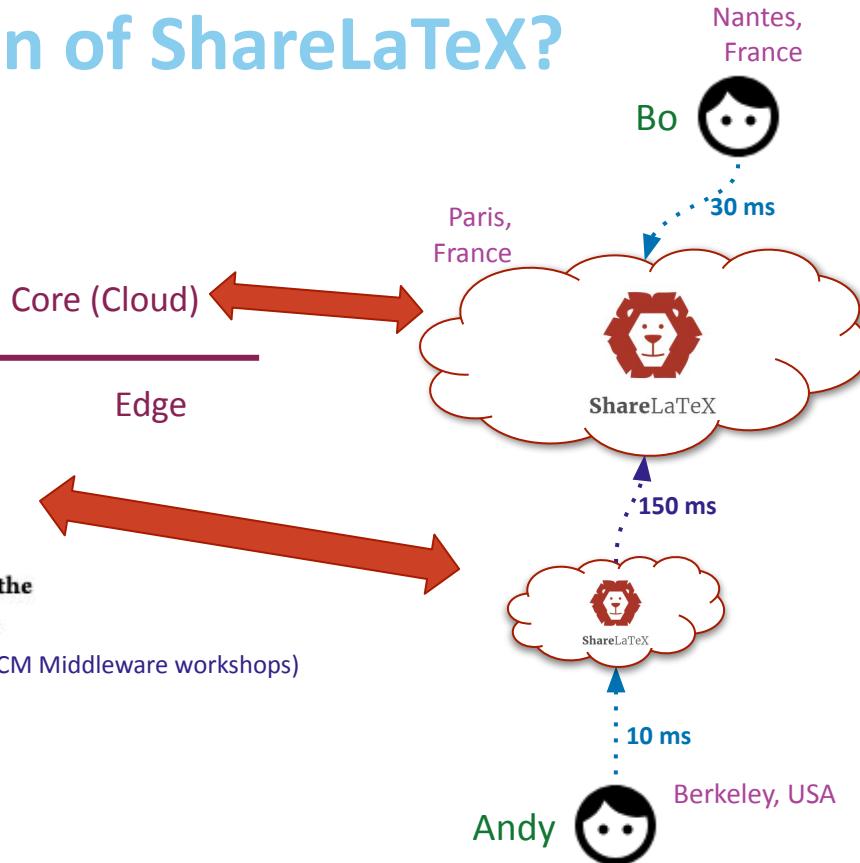
Source: Tato et al., *ShareLatex on the Edge*, ME'CC18 (ACM Middleware workshops)



Locality?



Network partitions?



Bases for an approach: locality is crucial



Local-first:

- Always able to serve local requests
- Minimize communications between DCs



Collaborative-then:

- Share resources across sites on demand
- Replicate them where needed

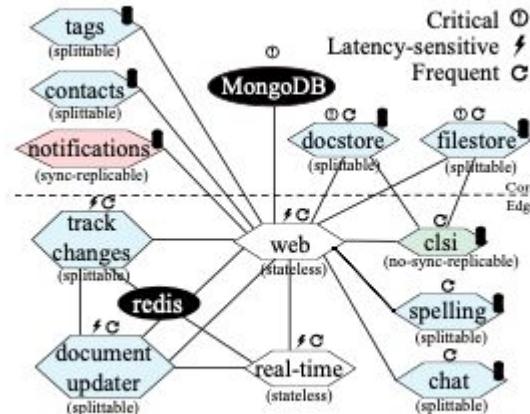


Figure 3: Our split and placement. Critical services in the core, latency-sensitive and frequent services at the edge.

Autonomous Cloud Applications: a first step

Autonomous instances: provides **locality** and **robustness**

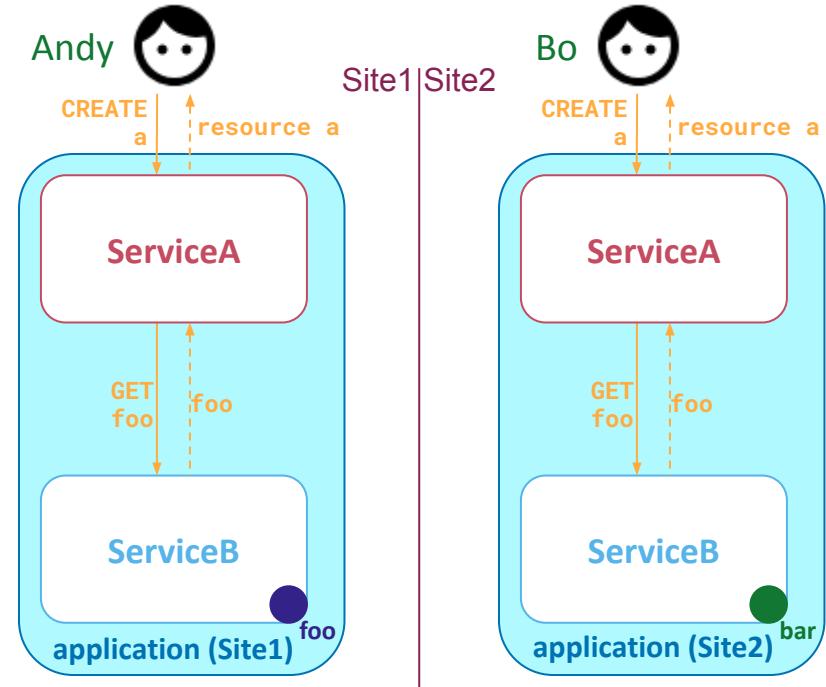
Andy and Bo use their own application, the closest:

Andy: application create resourceA
--sub-resourceB **foo**

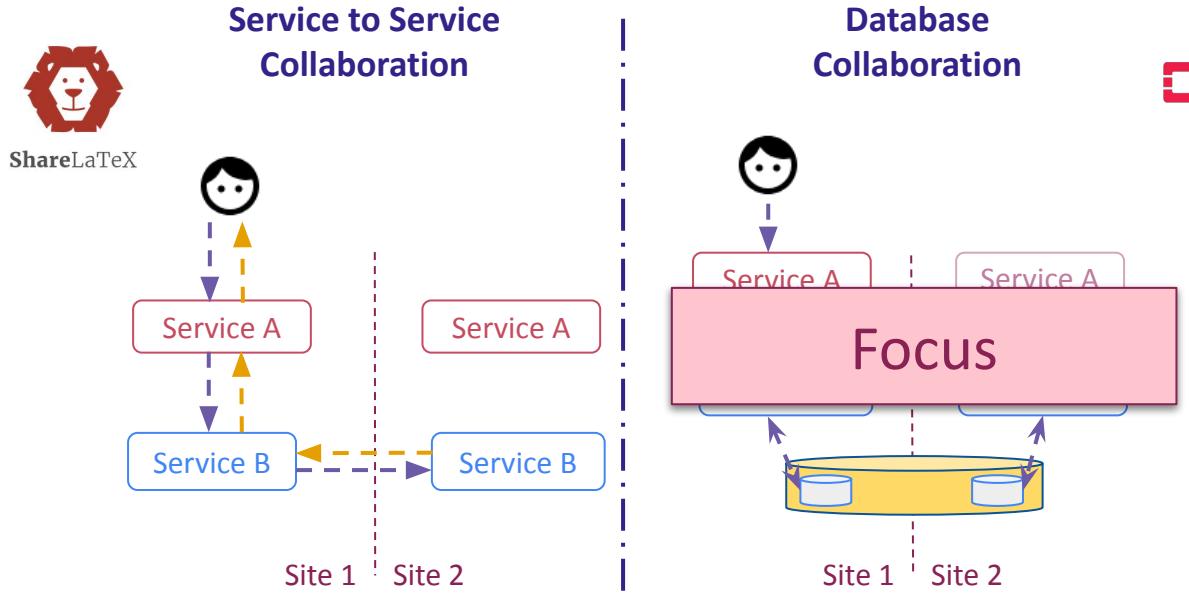
Bo: application create resourceA
--sub-resourceB **bar**

Without collaborations between instances,
what is the point?

Andy: application create resourceA
--sub-resourceB **bar** ?



Different approaches for collaborations

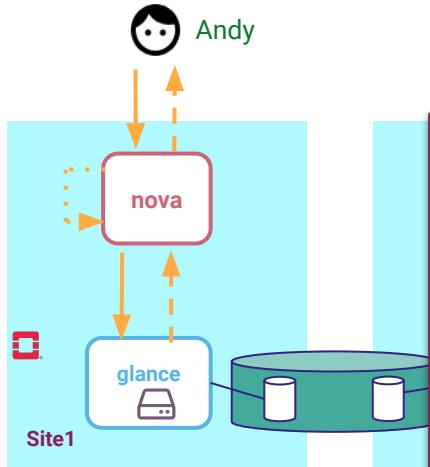


Problems:

The collaboration code is tangled in the business code

Resources have a context and a scope, sometimes side effects

Unanticipated issue



Coupling domain and collaboration code is not always an option:

- OpenStack is 13M LoC of Python (1.5M for core projects)[#]
- Core developers do not want to add complexity (geo-distribution is not that important for them)

```
- import os  
+ import urllib
```

```
@route('/image/{name}')  
def get_image(name: String) -> Blob:  
    # Get the path of the image stored
```

```
le:///path/debian.qcow  
ry(f  
path FROM images  
[S "{name}";'])  
oto://url/debian.qcow  
ch(name)
```

*the path on the disk to get
the image*

```
-     image = os.open(path, 'rb')  
+     # Follow the link to get the qcow image  
+     image = urllib.urlopen(path)  
     return image
```

Data sharing ≠ resource sharing

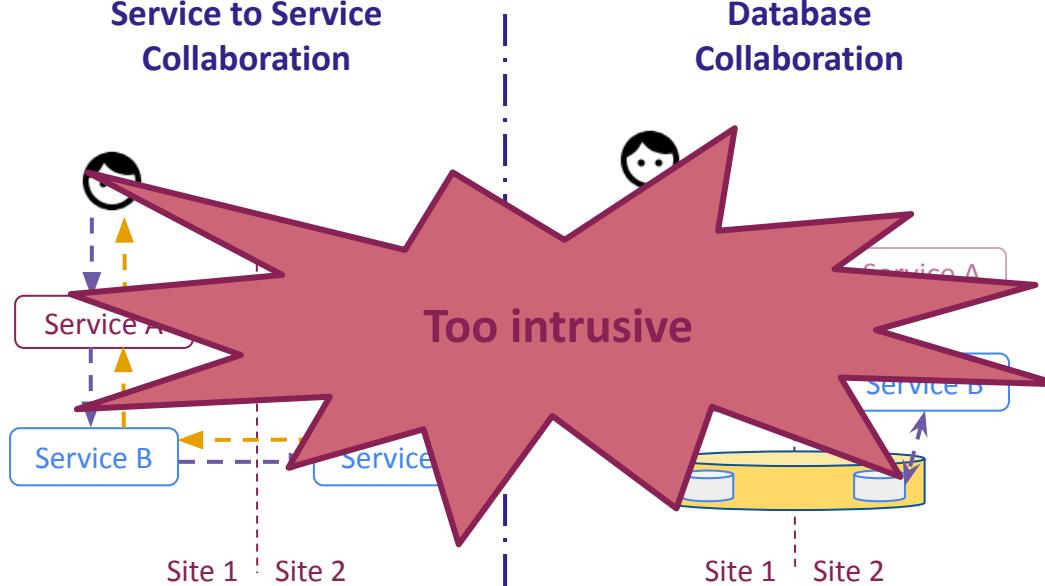
- A resource is data (+ effects)
- OpenStack VMs, nets, images, volumes **do not fit** in the DB
- App. code** has to **take into account** the distribution

[#] https://www.openhub.net/p/openstack/analyses/latest/languages_summary

Slide originally made by R.-A. Cherreau

Different approaches for collaborations

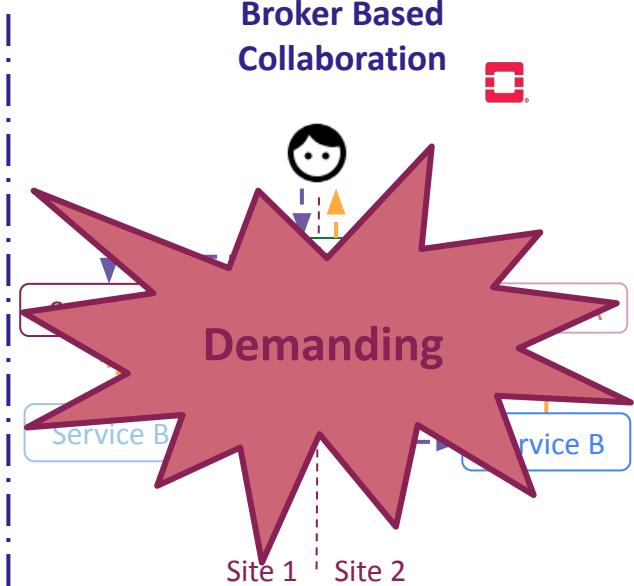
Service to Service Collaboration



Database Collaboration



Broker Based Collaboration



Problems:

The collaboration code is tangled in the business code

Resources have a context and a scope, sometimes side effects

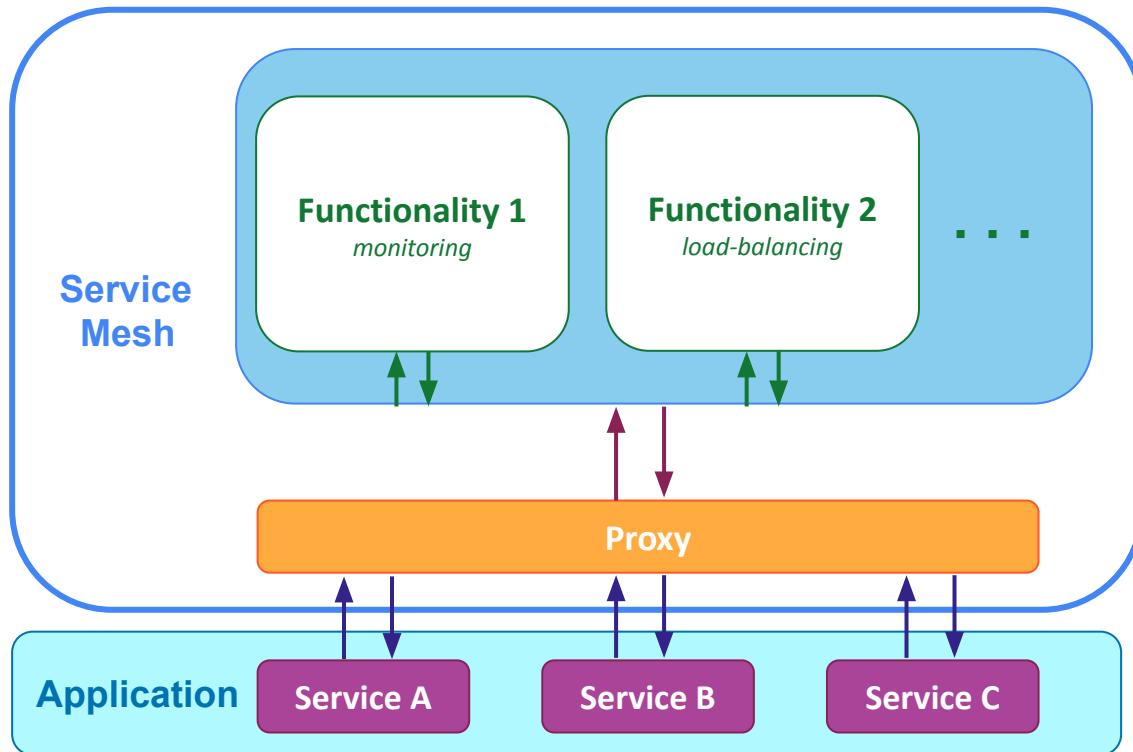
Brokers need to be developed for each service

The different approaches are either intrusive or tedious

Is it possible to find a **generic** solution to **externalize geo-distribution** concerns?

Lifecycle of service based applications can be managed externally through mechanisms such as a **service mesh**.

A simplified view of a service mesh



Purpose:

Manage some functionalities outside of the business code

Example of functionalities:

- load-balancing
- monitoring
- ...
- ...
- geo-distribution ?

Research questions, more specifically

-  Is it possible to use applications developed for the Cloud on Edge infrastructures without changing their code *leveraging a service mesh-like approach* ?

-  Use case: OpenStack, a Cloud infrastructure manager (and then Kubernetes)
 - OpenStack: complex application / 186 services (8 core services)
 - If OpenStack can be geo-distributed, we can propose a solution to manage an **Edge infrastructure** (which is an open challenge too)



State of the Art (cont)

Other academic and industrial solutions



Criteria

- Non-intrusive** no changing the code
 - Generic** should work on as many applications as possible
 - Local-first** ability of the solution to use only the local site
 - Collaborative-then** collaboration between sites
 - Network partition** robust to network partition
 - Dynamic** execution of requests handled dynamically
-
- On-demand** execution chosen per request by the users
 - Decentralized** decentralization of the solution
 - Peer-to-Peer** equality of privileges on the different sites

Comparison (more ☁ = more 😊)

	Approach	Non-intrusive	Genericity	Local-First	Collaborative-then	Resiliency to Network Partitions	Dynamicity	On-demand	Decentralized	P2P
Resource management on Edge Infrastructures	Decentralized CloudLightning	☁️☁️☁️	☁️☁️☁️	-	☁️☁️☁️	☁️	☁️☁️☁️	☁️	☁️☁️☁️	☁️☁️☁️
	Liqo	☁️☁️☁️	☁️☁️	☁️☁️☁️	☁️☁️☁️	☁️☁️☁️	☁️☁️	☁️☁️	☁️☁️☁️	☁️☁️
	HYDRA	☁️☁️☁️	☁️☁️	☁️	☁️☁️☁️	☁️☁️☁️	☁️☁️	☁️	☁️☁️☁️	☁️☁️☁️
	OneEdge	☁️☁️☁️	☁️☁️☁️	☁️☁️	☁️☁️☁️	☁️	☁️☁️☁️	☁️	☁️☁️	☁️☁️
	OpenYurt	☁️☁️☁️	☁️☁️	☁️☁️	☁️☁️	☁️	☁️☁️	☁️	☁️	☁️
	P2P OpenStack	☁️☁️☁️	☁️☁️	☁️☁️☁️	☁️☁️	☁️☁️	☁️☁️	☁️☁️	☁️☁️☁️	☁️☁️☁️
Service meshes	Istio	☁️☁️☁️	☁️☁️	☁️☁️☁️	☁️	☁️	☁️	☁️	-	-
	Linkerd	☁️☁️☁️	☁️☁️	☁️☁️☁️	☁️	☁️	☁️	☁️	☁️	-
Towards Edge-native applications	Scalable Gabriel	☁️	☁️☁️	☁️	☁️	☁️	-	-	☁️	☁️
	Colony	☁️	☁️☁️☁️	☁️☁️☁️	☁️☁️☁️	☁️☁️☁️	☁️☁️☁️	☁️☁️	☁️	☁️

The approach: Cheops and scope-lang

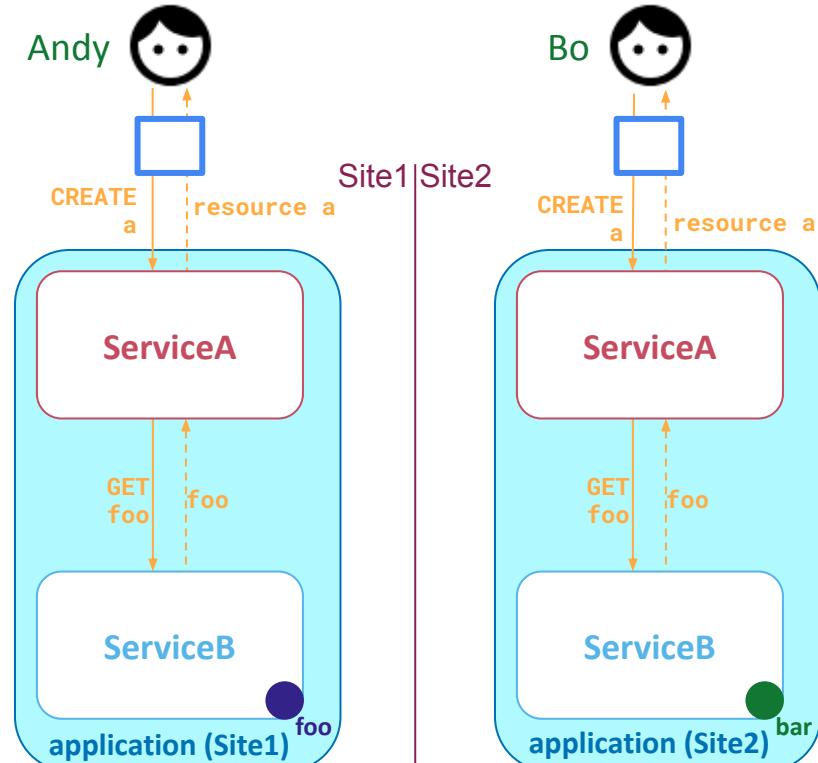
Autonomous Cloud Applications: a first step

Andy and Bo use their own application, the closest:

Andy: application create resourceA
--sub-resourceB **foo**

Bo: application create resourceA
--sub-resourceB **bar**

We need a solution to manage the
geo-distribution externally and in a generic
manner



Scope-lang

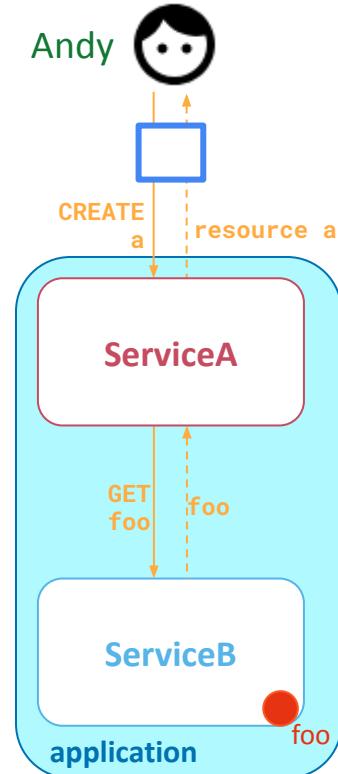
- Cloud icon A DSL to **manage resources & integrated with native CLI**
- Cloud icon A scope-lang expression contains:
 - Cloud icon **Collaboration & service** information
 - Cloud icon **Location** information
- Cloud icon Gives users a set of operations they can use to decide **where** and **how** a request will be executed.

For example:

```
application create resourceA --sub-resourceB foo
```

≡

```
application create resourceA --sub-resourceB foo  
--scope {ServiceA: Site1,  
         ServiceB: Site1}
```

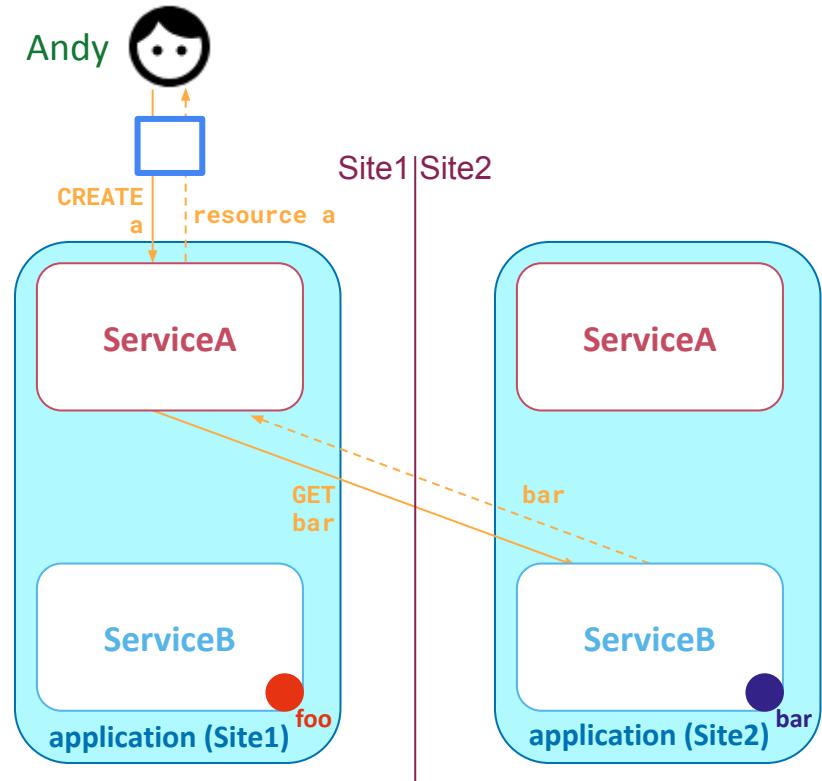


Scope-lang

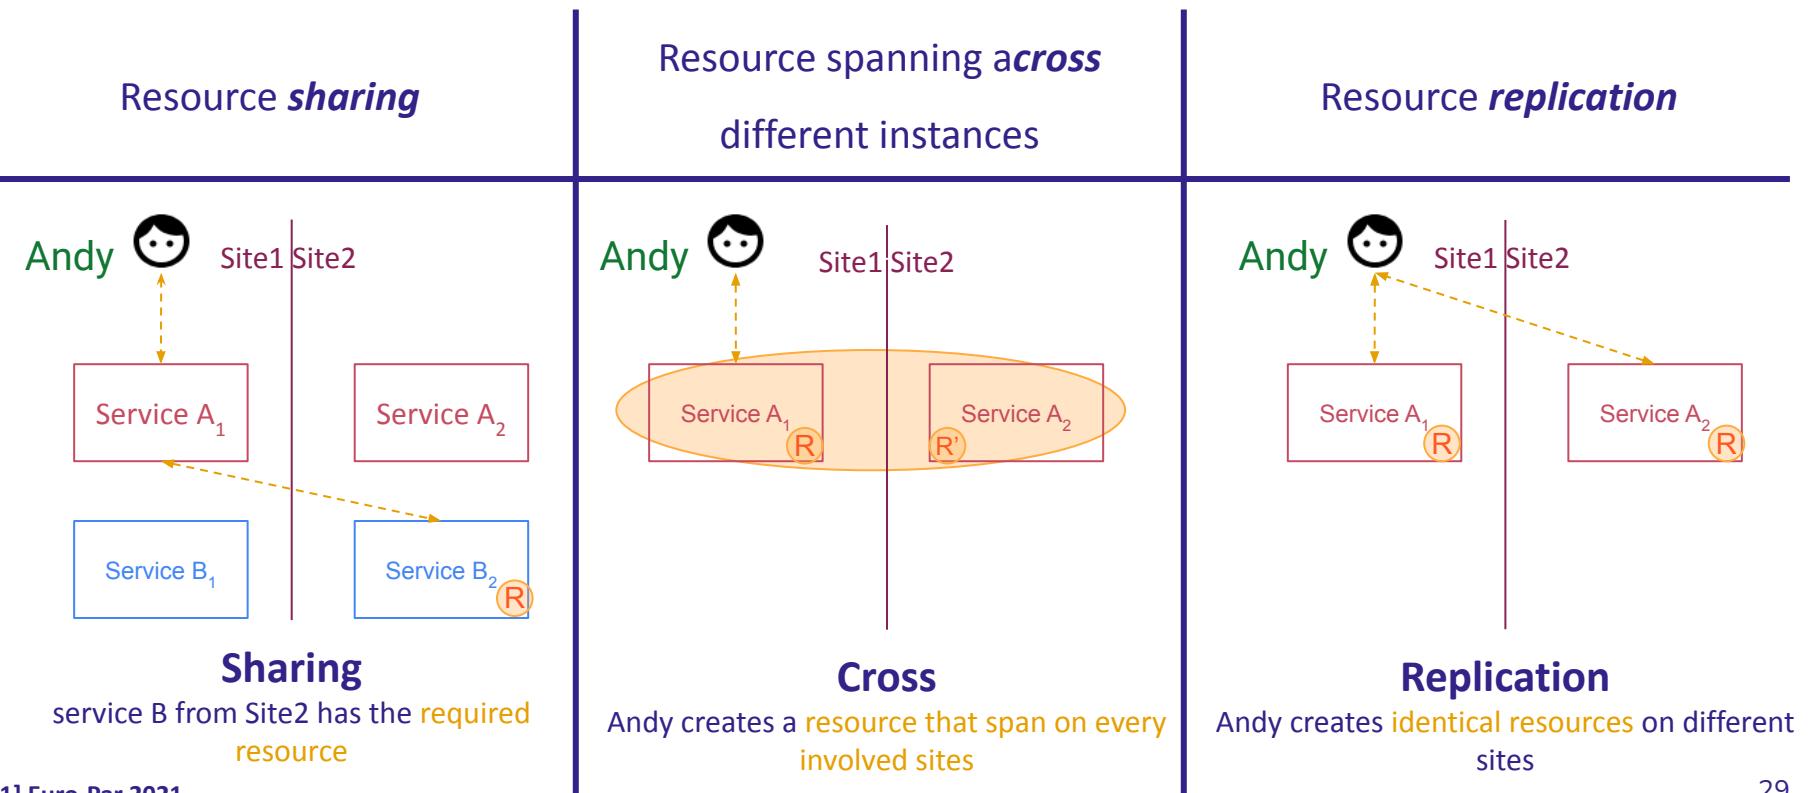
Scope-lang gives users a set of operations they can use to decide where and how a request will be executed.

For example:

```
application create resourceA
  --sub-resourceB bar
  --scope
    {ServiceA: Site1,
     ServiceB: Site2}
```



3 types of Collaborations^[1]

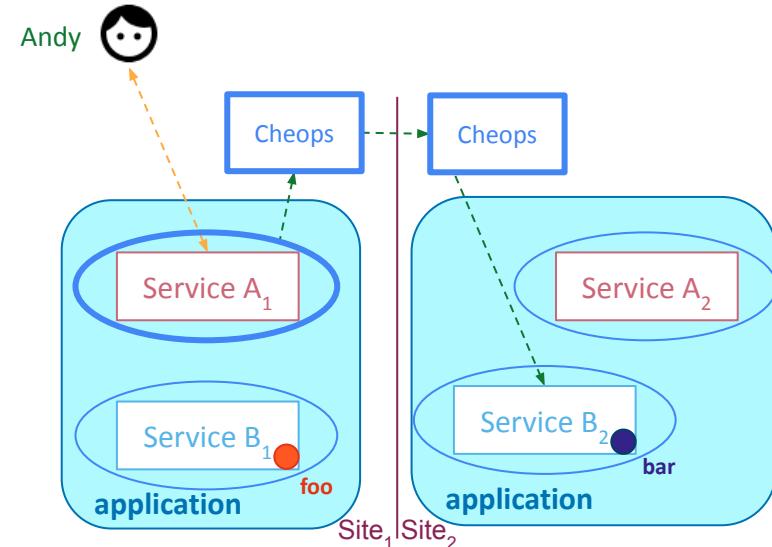


Sharing^[1]

Service B from Site2 has the **required resource (bar)**

Andy defines the **scope** of the request into the CLI. The **scope** specifies **where** the request applies.

```
application create resourceA  
    --sub-resourceB bar  
    --scope { serviceA: Site1 ,  
              serviceB: Site2 }
```



Sharing has been validated on OpenStack

```
openstack server create [...] myvm --scope {Nova: Berlin, Glance: Paris}
```

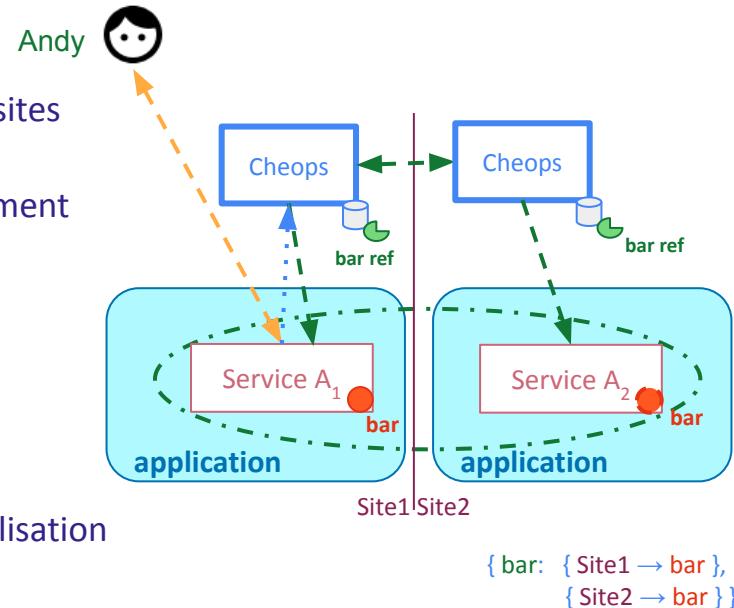
Cross^[1]

Cross creates the illusion of a **Single Resource** across the involved sites

Initial proposal made by Sarmiento et al.^[2], with intrusive development efforts made for Openstack

- Cloud icon: Cross maintains the availability of the resource
- Cloud icon: Cross is based on **Aggregation & Divisibility** principles

Currently still a work in progress by Geo Johns Antony for a generalisation



[1] ICSOC 2022

[2] Multi-site Connectivity for Edge Infrastructures : DIMINET: Distributed Module for Inter-site NETworking

Replication^[1,2]: a theory

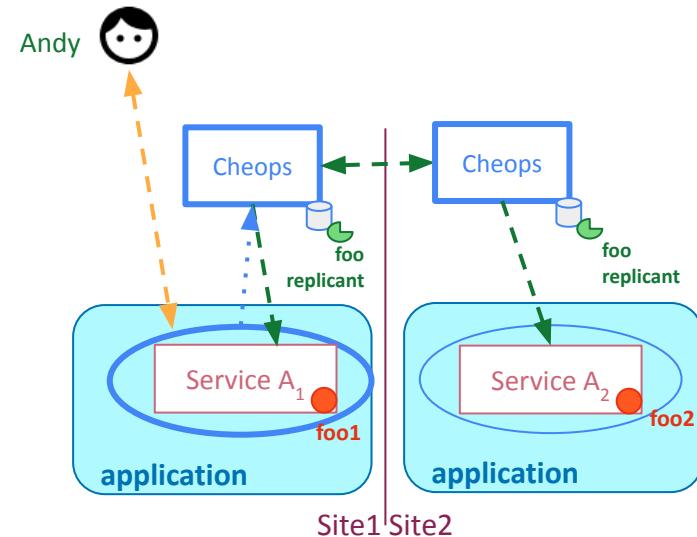
Andy defines in the scope **where the resource will be replicated**
The resource (managed by Service A) will be created on both sites

```
application create foo --scope { serviceA: Site1 & Site2 }
```

The goals of the replication are:

- cloud icon **lowering latency** when using a resource close
- cloud icon **increasing robustness** towards partition by allowing the use of the resource locally

Replicas information are stored in a **replicant**, associating the resource to **the location of the replicas and how to find them locally**.



[1] Euro-Par 2021

[2] XP 2021 Workshops

A view of the replica model: replicants



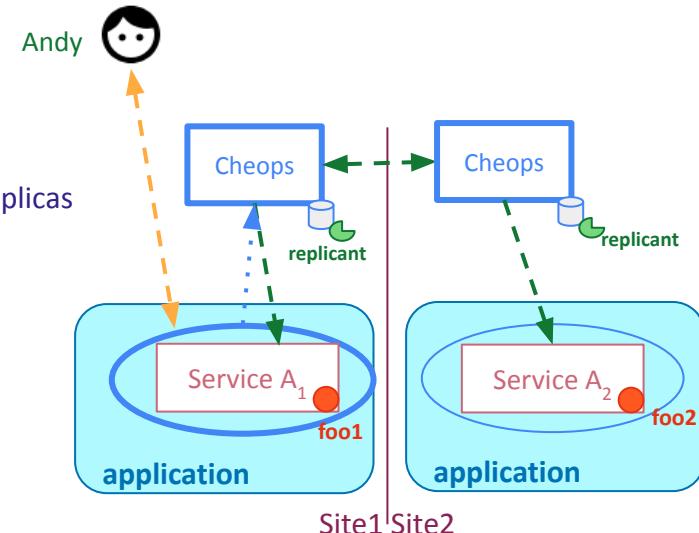
These replicants store:

- a meta ID (identifiant), generated by Cheops to keep track of the replicas
- each site where the replicas are located
- for each site, the local ID to find the replica when needed



Example:

- Andy wants two replicas of `foo` on Site1 and Site2
- Site1 creates the resource and assign `333` as local ID to this resource.
- Site2 does the same with the ID `111`.
- Cheops generates a meta ID: `42`
- The replicant is then: `{ 42: { Site1: 333, Site2: 111 } }`

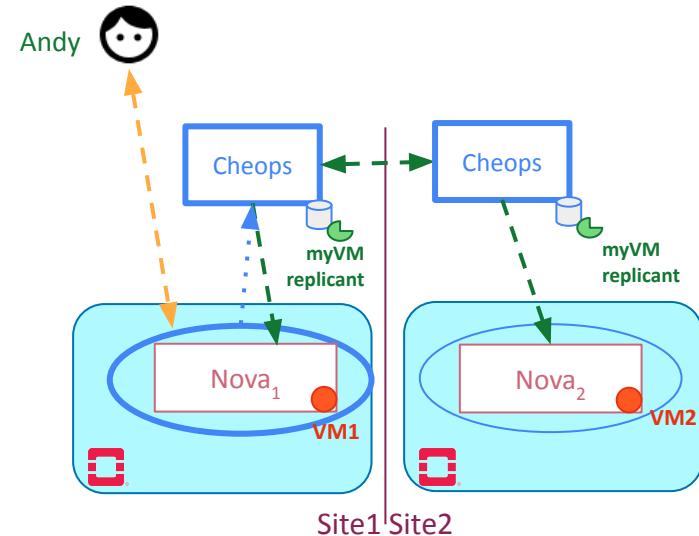


```
foo = { 42: { Site1 → 333,  
              { Site2 → 111 } } }
```

Replication, in the flesh

```
openstack server create myVM --scope { Nova: Site1 & Site2 }
```

- Replication follows an **eventual consistency** (Raft)
- Consistency is maintained **through the API** (at the CRUD level, resources are considered as black boxes)



Like for sharing, resources are considered as black boxes for **genericity**

The premises for an implementation of Cheops

Our specification/requirements for the execution:

- ☁ Follows the **requirements** defined in the State-of-the-Art
 - ☁ **Modular**
 - ☁ Ability to integrate **scope-lang** and the **collaborations**
- 
- ☁ Non-intrusive
 - ☁ Generic
 - ☁ Local-first
 - ☁ Collaborative-then
 - ☁ Network partitions tolerance
 - ☁ Dynamic
 - ☁ On-demand
 - ☁ Decentralized and P2P

Primary targets: OpenStack and Kubernetes

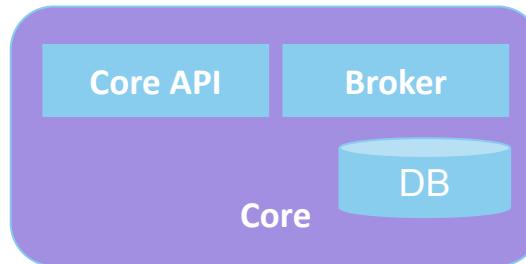


Our solution: Cheops

Cheops divided into 2 major modules:

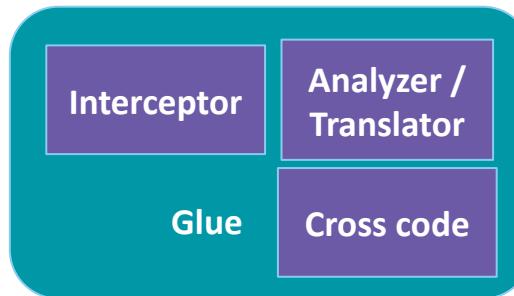
- **Cheops Core:**

- Management layer responsible for the deployed application metadata & P2P interactions
- Generic across applications

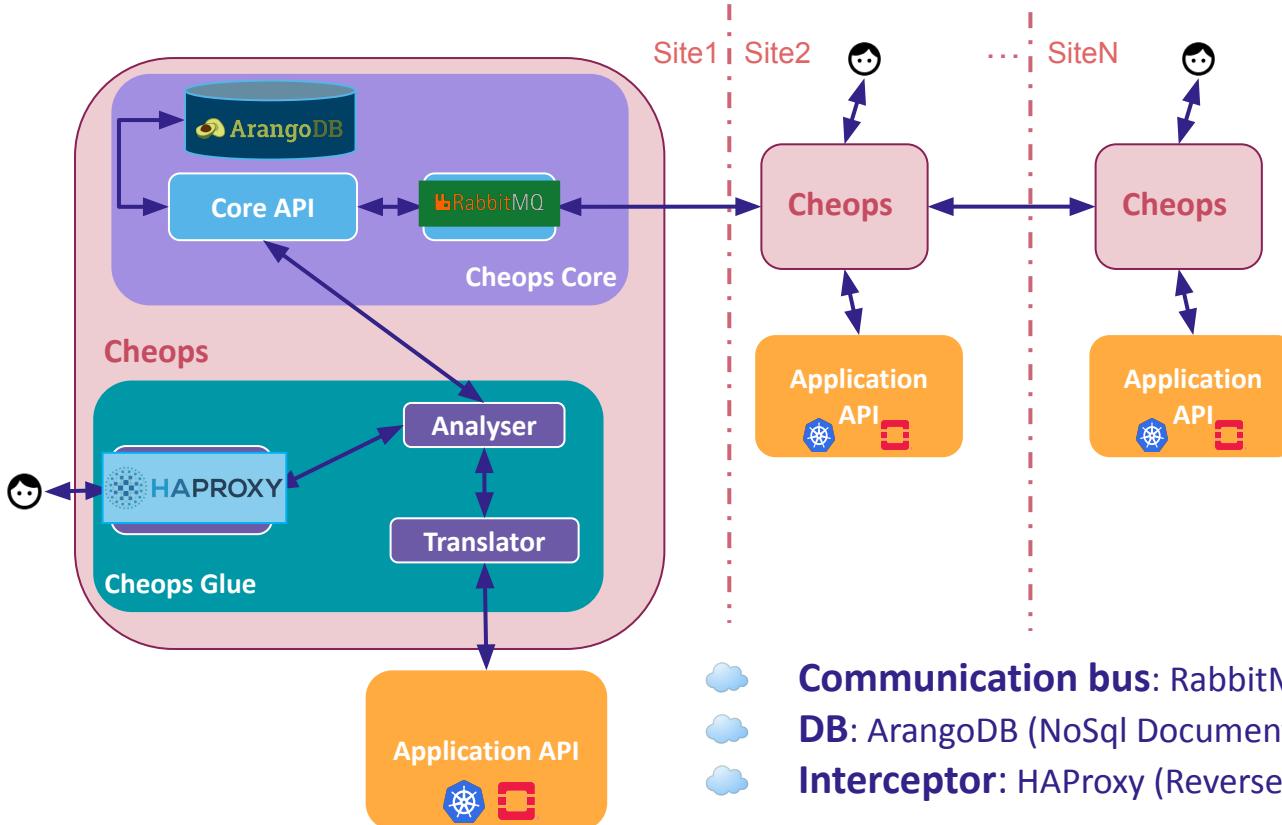


- **Cheops Glue:**

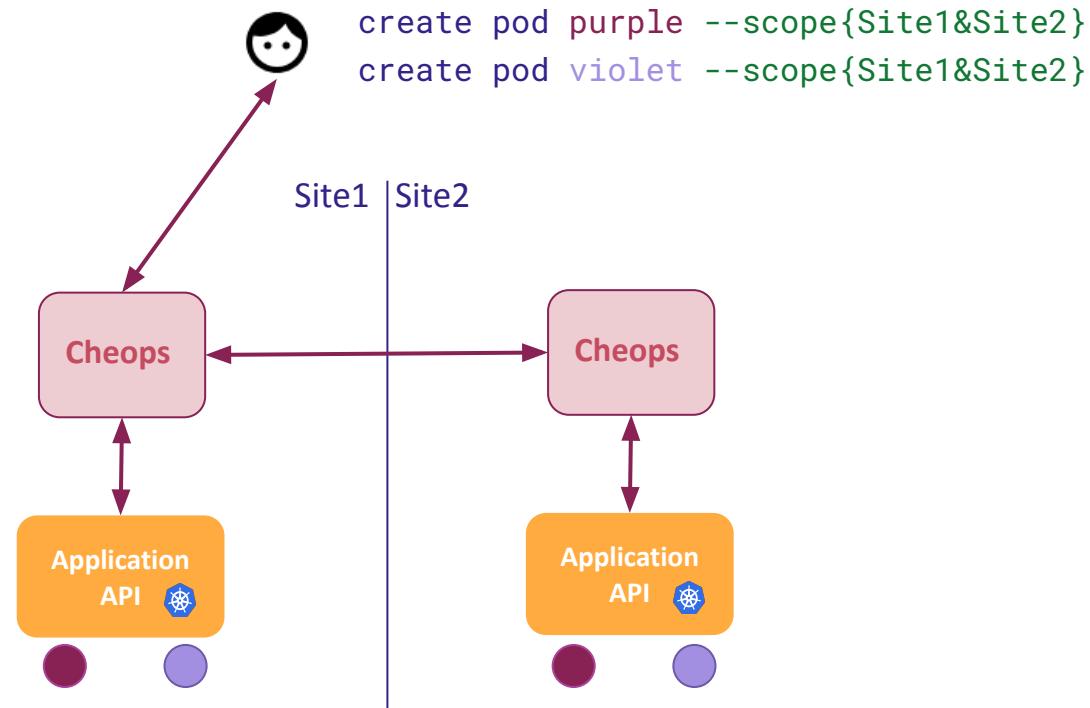
- Intermediate layer responsible for translating interactions with the application
- Dependent to an application



Cheops architecture



Sandbox experiments on Kubernetes



Note: Sharing does not make sense on Kubernetes because it only has one service

Sandbox experiments on Kubernetes

Operation	Location	Result
kubectl create pod purple --scope{Site1&Site2}	Site1	Pod purple created on Site1 and Site2
kubectl create pod violet --scope{Site1&Site2}	Site2	Pod violet created on Site1 and Site2
	kubectl get pod violet	Site1 Pod violet from Site1 is displayed
	kubectl get pod violet	Site2 Pod violet from Site2 is displayed

Note: ongoing experimentations (scalability, network partitions, ...)

Resource dependencies Model

A Generic Pattern

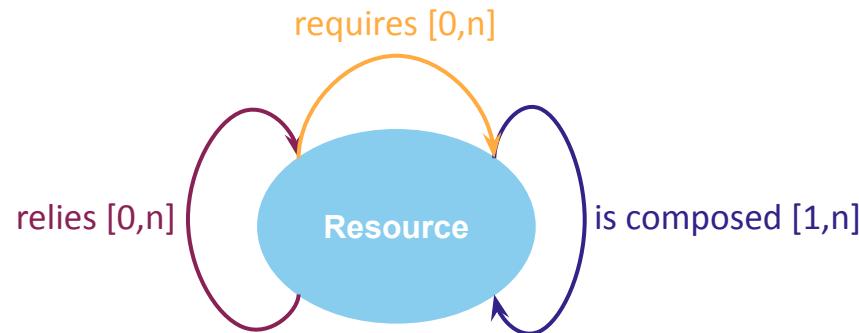
Services relies on resources like VMs, pods, services, networks, etc. that may be **interdependent**.

To geo-distribute applications in a non-intrusive manner:

- ☁️ Involves handling resources dependencies between these micro-services
- ☁️ This dependencies handling must be also generic

Can we find a generic pattern from multiple applications to manage these **dependencies**?

Patterns in dependencies^[1]



Example:

A vm **requires** an image

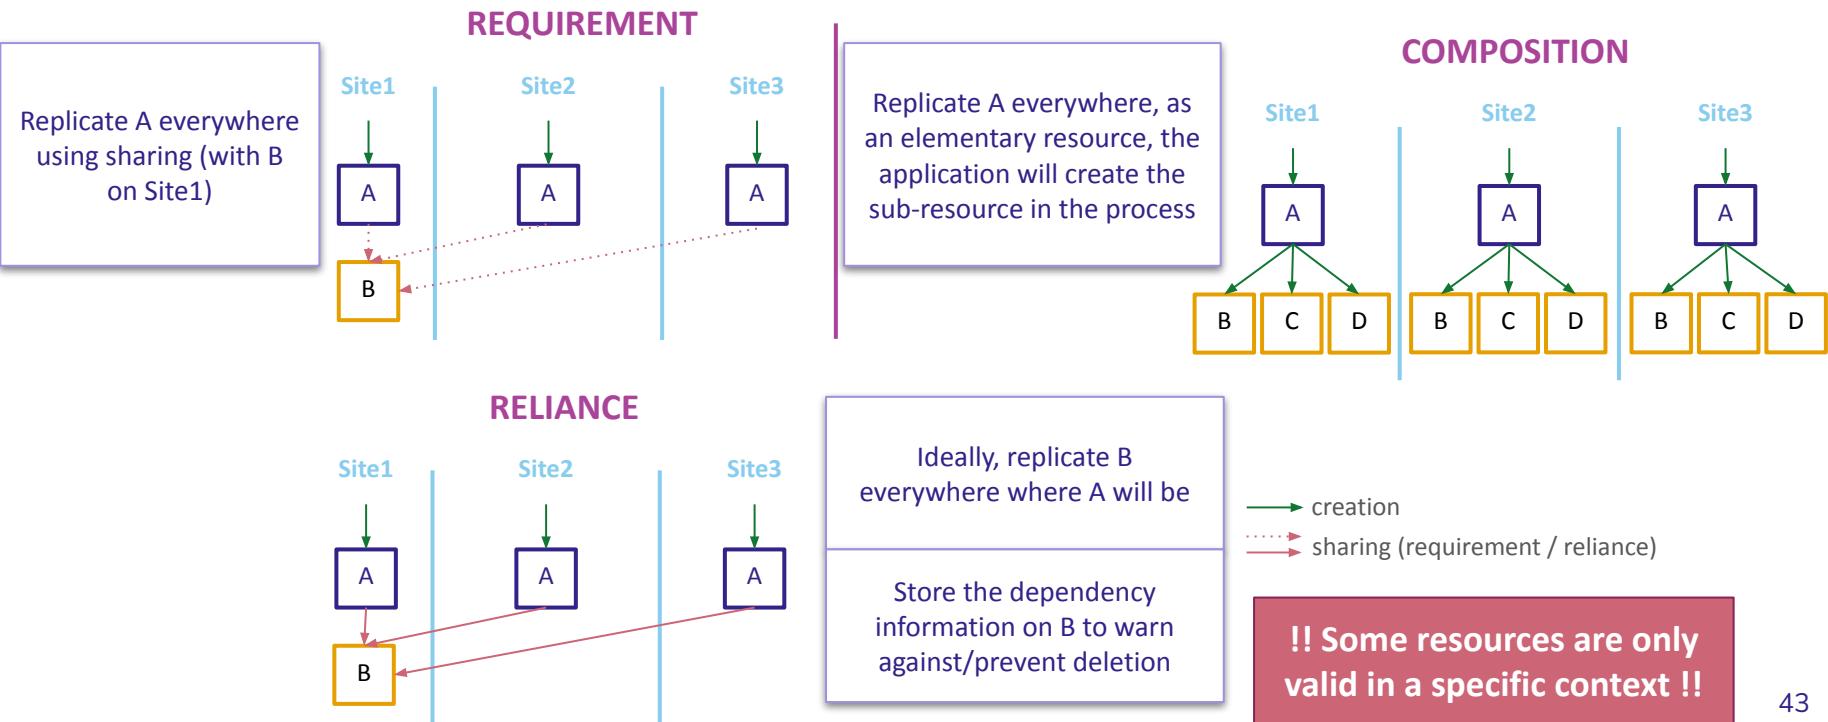
Example:

A vm **relied** on an IP

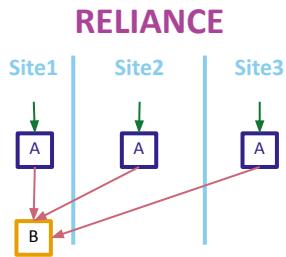
Example:

A heat stack **is composed of** vms

How to manage dependencies - Replication



Invalid collaborations



Ideally, replicate B everywhere where A will be

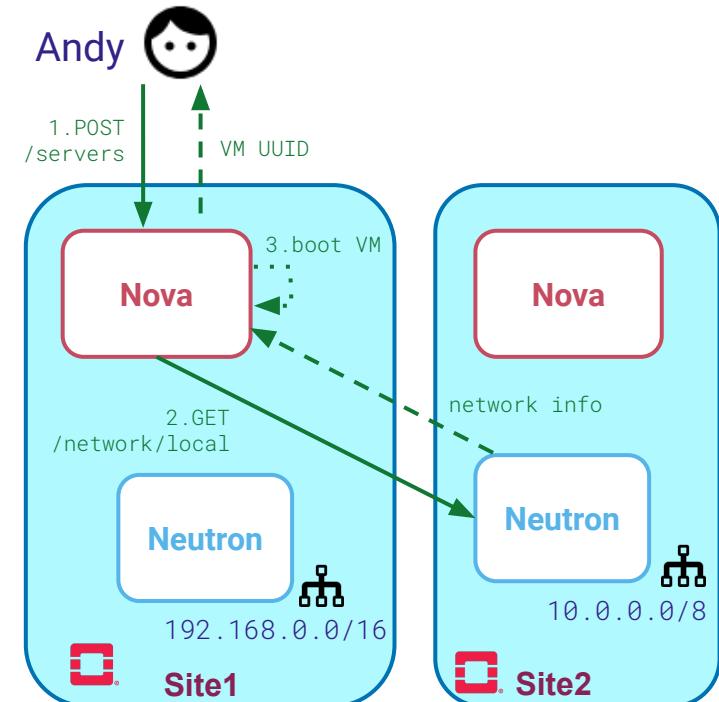
Store the dependency information on B to warn against/prevent deletion

!! A vm relies on an IP !!

```
openstack server create myVM --ip myIPSite2 ??  
--scope { Nova: Site1 & Site2,  
Neutron: Site2 }
```

- Wrong net info for the compute service (e.g., 10.0.0.2 instead of 192.168.0.2)
 - :(the VM on Site1 is unreachable

Resources are valid only in a specific context



Boot VM on Site1 with a network from Site2

Conclusion and Perspectives

Conclusion on scope-lang and Cheops

Cheops/**scope-lang** is a Proof of concept to geo-distribute Cloud applications at the Edge. It is:

- ☁ **Non-intrusive** no change has been made to the application code
- ☁ **Generic** to applications made of services
- ☁ **Local-first** autonomous instances of the application everywhere
- ☁ **Collaborative-then** collaborations thanks to scope-lang
- ☁ **Network partitions tolerance** local-first and collaborations allows some tolerance
- ☁ **Dynamic** scope-lang applies to any fine-grained request
- ☁ **On-demand** scope-lang specifies only location on-demand, otherwise, the request is local
- ☁ **Decentralized and P2P** fully decentralized with no central authority or hierarchy

Cheops: ongoing/future work

- ☁ More **experimentations** on partitioning, scaling, tests on OpenStack
- ☁ **Dependencies management** to be formalised and **invalid collaborations** prevented with a dependency tree
- ☁ More **collaborations** (with Cross under way)

Collaborations and Scope-lang

Cloud icon **Or**

- Cloud icon `application resourceA create --scope { serviceA: Site1 || Site2 }`

Cloud icon **Around**

- Cloud icon `application resourceA create --scope { around: Site1, 50ms }`

Cloud icon **&set (or replicaset, inspired by Kubernetes)**

- Cloud icon `application resourceA create --scope { set: 2@Site1, 3@Site2 }`

- Cloud icon managed by an autonomous control loop

Research question

Is it possible to use applications developed for the Cloud on Edge infrastructures without changing their code?

Through this Phd Defense I hope I demonstrated that it is indeed possible 😊

Publications

Conferences

- **ICSO2022** Delavergne, M., Antony, G.J., Lebre, A. (2022). Cheops, a Service to Blow Away Cloud Applications to the Edge. In: Troya, J., Medjahed, B., Piattini, M., Yao, L., Fernández, P., Ruiz-Cortés, A. (eds) Service-Oriented Computing. ICSOC 2022. Lecture Notes in Computer Science, vol 13740. Springer, Cham., doi: 10.1007/978-3-031-20984-0_37.
- **Euro-Par 2021** Cherrueau, RA., Delavergne, M., Lebre, A. (2021). Geo-distribute Cloud Applications at the Edge. In: Sousa, L., Roma, N., Tomás, P. (eds) Euro-Par 2021: Parallel Processing. Euro-Par 2021. Lecture Notes in Computer Science(), vol 12820. Springer, Cham., doi: 10.1007/978-3-030-85665-6_19.



Journal

- **IEEE TPDS 2022** R.-A. Cherrueau et al., "EnosLib: A Library for Experiment-Driven Research in Distributed Computing", in IEEE Transactions on Parallel and Distributed Systems, vol. 33, no. 6, pp. 1464-1477, 1 June 2022, doi: 10.1109/TPDS.2021.3111159.

Workshop

- **XP 2021 Workshops** Delavergne, M., Cherrueau, RA., Lebre, A. (2021). A Service Mesh for Collaboration Between Geo-Distributed Services: The Replication Case. In: Gregory, P., Kruchten, P. (eds) Agile Processes in Software Engineering and Extreme Programming – Workshops. XP 2021. Lecture Notes in Business Information Processing, vol 426. Springer, Cham. doi: 10.1007/978-3-030-88583-0_17.125



Research reports

- **RR 2022** Marie Delavergne, Geo Johns Antony, Adrien Lebre. Cheops, a service to blow away Cloud applications to the Edge. [Research Report] RR-9486, Inria Rennes - Bretagne Atlantique. 2022, pp.1-16. <hal-03770492v2>.
- **RR 2020** Ronan-Alexandre Cherrueau, Marie Delavergne, Adrien Lebre, Javier Rojas Balderrama, Matthieu Simonin. Edge Computing Resource Management System: Two Years Later!. [Research Report] RR-9336, Inria Rennes Bretagne Atlantique. 2020. <hal-02527366v2>.

Outside of the scientific community

- **OpenInfra Summit 2022** Geo Johns Antony, Marie Delavergne and Baptiste Jonglez. Cheops - Can a "service mesh" be the right solution for the Edge?, <https://www.youtube.com/watch?v=7EZ63DMRJhc>, OpenInfra Summit 2022, Berlin

